

Thinker AI

V1 Architecture Specification

Complete Edition — June 2026

Synthesized from Founder Constitution (Extended), Constitution v1.0, and Founding Charter v1.0. This complete edition incorporates all original sections (1–17) together with four architectural addenda: Section 18 (Paid Plan Architecture), Section 19 (Resilience Architecture), Section 20 (Dynamic Routing Engine), and Section 21 (Think Credits & Thinking Levels).

Table of Contents

1. Purpose of This Document
2. V1 Scope Statement
3. Core Philosophy → Implementation Mapping
4. System Architecture ★ REVISED
5. Agent Specifications ★ REVISED
6. Model & Cost Strategy ★ REVISED
7. Data Flow & State ★ REVISED
8. Agent Fallibility & Error Containment ★ REVISED
9. Pipeline Observability & Post-Delivery Feedback
10. Technology Stack
11. Build Order ★ REVISED
12. V1 Launch Checklist
13. Explicitly Out of Scope for V1
14. Open Decisions for Founder Review
15. Goal Intelligence System
16. Output Pipeline — Preview Before Build
17. Multi-Language Support
18. Paid Plan Architecture ★ NEW
19. Resilience Architecture — Continuous Delivery Guarantee ★ NEW
20. Dynamic Routing Engine — Implementation Specification ★ NEW
21. Think Credits, Pricing & Thinking Levels ★ NEW

1. Purpose of This Document

The three founding constitutions describe Thinker AI's vision and philosophy: a multi-agent system that understands goals, plans work, debates solutions, and delivers verified outcomes. They do not specify how to build it.

This document translates that philosophy into a concrete, buildable Version 1 (V1) architecture. It defines every agent, the data that flows between them, the technology choices, and the order of implementation — scoped deliberately small so V1 can actually ship.

Guiding constraint: all nine agent roles named across the three constitutions are preserved, but each is implemented with the minimum logic needed to fulfil its role. Complexity is added in later versions, not V1. Three additional agents are introduced: Design Agent (Section 5.10), Strategy Agent (Section 5.12), and the expanded Decision Memory system (Section 7.3) — bringing the total to twelve agent roles in V1.

2. V1 Scope Statement

2.1 What V1 Is

A web-based chat application where a user describes a goal (build an app, a website, a game, or a general task), and a dynamic pipeline of twelve lightweight agents — coordinated by Thinker Core — clarifies the request, thinks strategically about it, plans it, builds a single deliverable (including basic image/graphic assets where needed), reviews it, and returns one final answer to the user. Thinker AI supports 100+ languages and responds automatically in the user's language throughout every stage of the pipeline (Section 17).

2.2 What V1 Is Not

- Not a fully autonomous multi-day project system (that is a later phase).
- Not a mobile app (V1 ships as a responsive web app).
- Not running every request through full multi-model debate — consensus is invoked selectively (Section 6.4).
- Not retaining credentials, passwords, or payment data in memory.

2.3 V1 Success Criteria

#	Criterion
1	User can chat normally and get direct answers without triggering the full agent pipeline
2	User can request a small project and the system asks clarifying questions before building
3	All twelve agent roles are present in code, even if some are simple rule-based stubs
4	Every final deliverable passes through Reviewer before being shown to the user
5	Multi-model consensus works for at least one agent (Judge) end-to-end

#	Criterion
6	Project memory persists within a session and reloads on return
7	For projects with a visual asset, the user explicitly approves the final image
8	Smart Clarification correctly detects vague goals and activates 3-Level Deep mode
9	Signature Question fires on every project-type request before Builder runs

3. Core Philosophy → Implementation Mapping

Constitution Principle	V1 Implementation
Understand first, plan second, execute third, verify last	Fixed pipeline order, enforced in code — but dynamic routing on failure (Section 4.1)
Requirement Clarity Law — ask minimum questions before large work	Clarification Agent has a hard gate; Builder cannot run until requirements are marked complete
Token Economy / Resource Protection Law	Cheap single-model calls by default; multi-model consensus only for high-stakes decisions
One Goal. One Team. One Final Answer.	Thinker Core merges all agent outputs into exactly one response object before returning
No single model/agent/opinion is final — verified consensus is required	Consensus Agent approval is required before Thinker Core marks a deliverable 'final'

The Founder Law

Thinker AI must not immediately build what the user asks for. It must first determine whether the user's goal, problem, and assumptions are valid. This is the single most important rule that separates Thinker AI from every other AI builder tool. Building first and validating later wastes the user's time, money, and trust. This law is enforced architecturally: the pipeline order makes it structurally impossible to reach Builder without first passing through Clarification Agent and Strategy Agent.

The Preview Before Build Law

Thinker AI must show the user what it intends to build before generating the final deliverable. User approval is required at each stage before the next stage begins. A ZIP file that surprises the user is a failure, even if the code is correct. See Section 16 for the full 5-stage output pipeline that implements this law.

4. System Architecture ★ REVISED

4.1 Pipeline Overview ★ REVISED

The pipeline no longer runs in a fixed 1→12 sequence. Thinker Core operates as a **dynamic routing engine**: after every agent completes, Thinker Core reads the agent's **next_action** output field and decides the next step. The twelve agent roles are unchanged; what changes is that Thinker Core now reacts to failure and uncertainty dynamically rather than propagating a broken state silently down a fixed chain.

Condition	next_action	Thinker Core Response
Simple chat / general question	direct_answer	Skip pipeline → single LLM call → return
Request is vague or incomplete	clarify	Halt pipeline → return questions to user
Requirements complete, project proposed	proceed	Run full pipeline in order
Builder output rejected by Reviewer	retry	Re-run Builder (same plan, same step)
Reviewer + Critic both flag deep issue	explain	Return to Planner with failure context
Planner→Builder loop exceeds 2 cycles	clarify	Return to Clarification for new requirements
Judge score borderline	escalate	Invoke Consensus Agent
All models for an agent unavailable	failover	Model Pool Manager assigns next provider pool
Any agent exceeds loop_count limit	halt	Pause pipeline, notify user, save state

Standard full-pipeline order (when next_action is 'proceed' throughout): Intent → Clarification → Strategy → Planner → Research → Builder → Reviewer → Critic → Judge → [Consensus if borderline] → Return to user.

4.2 Thinker Core Responsibilities ★ REVISED

Thinker Core is not itself an LLM call — it is the orchestration layer (plain backend logic). Original responsibilities retained, plus the following additions:

- Owns the pipeline state machine and dynamic routing table for each conversation.
- Decides whether a message needs the full pipeline or just a direct chat reply.
- Routes each plan step to Builder or Design Agent based on output type (Section 4.3).
- Passes structured data between agents (see Section 7) and enforces the Clarification gate.
- Merges all agent outputs into one final response and writes to memory after each stage.
- **Dynamic next-step decision:** reads next_action from every agent output before passing control forward.
- **Model Pool Manager:** maintains a pool of available models per agent tier. If a model call fails or times out (5-second threshold), silently reassigns to the next available model in the pool.
- **Loop detection:** tracks loop_counts per agent. If any agent is retried more than 3 times, or if the Planner→Builder→Reviewer cycle runs more than 2 full times, halts the pipeline and notifies the user.

- **Peer audit trigger:** when a failover model takes over a partially completed task, instructs it to audit what the previous model produced before continuing.
- **State preservation on halt:** on any halt, writes the full Pipeline State to the database before stopping. Resume is always possible from the last completed agent step.

4.3 Builder vs. Design Agent Routing

A single project can need both. Routing is decided per step, not per project:

Planner Step Output Type	Routed To
code, text, config	Builder Agent
image	Design Agent

If a step's output type is ambiguous, Thinker Core defaults to Builder and lets Builder request an image sub-step from Design Agent if genuinely needed.

4.4 Domain Picker — A Context Layer Above the Pipeline

Thinker AI's UI offers named domain agents (Music, DevOps, Security, Coding, etc.) that the user can pick before sending a message. The picker is a domain hint that flows into the one pipeline — it is input, not an alternate path. The twelve agents in Section 5 are the only system that ever produces a deliverable.

Agent	How the domain hint is used
Intent Agent	Skipped entirely when a domain is already selected — intent type is derived directly from the domain.
Clarification Agent	Uses domain-specific required fields where they exist.
Planner Agent	Receives the domain name in its prompt so step descriptions are framed appropriately.
Builder Agent	Receives the domain as additional system-prompt context, biasing toward domain-appropriate conventions.
Reviewer / Critic / Judge	Unaffected — quality and completeness checks apply the same way regardless of domain.

Implementation gap (current): the domain picker exists in the mobile UI but the selected domain is not yet included in the request sent to /api/pipeline. Closing this gap is Phase 10 in the build order (Section 11).

5. Agent Specifications ★ REVISED

All agents now include one additional output field read exclusively by Thinker Core for routing: **next_action** (proceed | retry | replan | clarify | escalate) and **reason** (string). Agents recommend next_action but do not execute it — Thinker Core is the sole actor on this field.

5.1 Intent Agent

Role: Classify the incoming message as chat, app request, website request, game request, or task request. Also outputs a best-guess domain label and the Thinking Level selection (Section 21.3). Skipped entirely when user has already selected a domain via the Domain Picker.

Input: Raw user message plus last five messages of context.

Output: Intent type, confidence score, domain label, Thinking Level.

V1 Implementation: Single LLM call with a classification prompt. No fine-tuning.

Model Tier: Fast/cheap.

5.2 Clarification Agent ★ UPDATED

Role: Enforce the Requirement Clarity Law AND run Smart Clarification — detect emotional context, surface hidden goals, challenge assumptions, and validate the idea before any technical work begins. This is the most critical agent: Wrong Goal → Wrong Strategy → Wrong Plan → Wrong Output.

Input: Intent type, full conversation, emotional tone signals.

Output: Completeness flag, smart questions, known requirements, assumption flags, Signature Question response.

V1 Implementation: LLM call against Smart Clarification Layer + 3-Level Deep Clarification + Constraint & Assumption Discovery + Signature Question Protocol.

Model Tier: Fast/cheap with enhanced prompt engineering.

5.3 Planner Agent

Role: Convert a clarified goal into an ordered list of concrete steps, using the Strategy Agent brief as context.

Input: Known requirements object + Strategy Agent brief + Signature Question response + Constraint findings.

Output: List of steps, each with id, description, output type, and dependencies.

V1 Implementation: Single LLM call producing a structured step list. No re-planning loop in V1.

Model Tier: Mid-tier reasoning.

5.4 Research Agent

Role: Gather external information when a step requires it. Optional — only invoked if Planner flags a step.

Input: A single step from the Planner's list, flagged as needing research.

Output: Findings and sources.

V1 Implementation: Optional. Uses web search. Skipped entirely for purely generative steps.

Model Tier: Fast/cheap with search tool access.

5.5 Builder Agent

Role: Produce the actual deliverable — code, copy, configuration, or content — for each Planner step. Most token-heavy agent.

Input: Step description, research findings (if any), relevant project memory.

Output: Artifact type, content, generated files.

V1 Implementation: Runs once per step, sequentially in V1. Parallel execution is a later optimisation.

Model Tier: Strongest available coding/reasoning model.

5.6 Reviewer Agent

Role: Run the quality control checklist: accuracy, logic, security, completeness, user-goal alignment. For images: checks relevance and appropriateness rather than the code checklist.

Input: Builder output plus original requirements.

Output: Pass/fail flag and list of issues with severity.

V1 Implementation: Single LLM pass against the five-point checklist. If it fails with high-severity issues, the request goes back to Builder once — capped at one retry in V1.

Model Tier: Mid-tier reasoning.

5.7 Critic Agent

Role: Independently look for weaknesses Reviewer's checklist-based pass might miss — edge cases, unclear UX, fragile assumptions. Critic's prompt explicitly withholds Reviewer's verdict until after Critic has formed its own opinion.

Input: Builder output plus Reviewer findings.

Output: List of concerns with an overall severity rating.

V1 Implementation: Single LLM call with adversarial/skeptical framing — deliberately different persona from Reviewer.

Model Tier: Mid-tier reasoning.

5.8 Judge Agent

Role: Score the deliverable on accuracy, feasibility, safety, cost, and completeness, then decide final approval. For high-risk categories (auth, payment, data deletion, permissions), Consensus fires automatically regardless of score.

Input: Builder output, Reviewer result, Critic concerns.

Output: Score per criterion and approval decision.

V1 Implementation: If clearly high or clearly low, Judge decides alone. If borderline, invokes Consensus Agent.

Model Tier: Strongest available reasoning model.

5.9 Consensus Agent

Role: Multi-model consensus used only when Judge needs a tie-breaker or for mandatory high-risk categories. Three personas: Conservative Reviewer, Pragmatic Engineer, Critical Skeptic — never provider names.

Input: The same question Judge was uncertain about.

Output: Each model's vote and reasoning, plus a final verdict.

V1 Implementation: Sends identical question to three different providers in parallel. Majority vote wins; ties default to conservative outcome.

Model Tier: Three different providers, called in parallel.

5.10 Design Agent

Role: Handle visual/graphic output — logos, illustrations, UI mockup imagery, posters, icons. Does not attempt 3D models or animation/video (see Section 13).

Input: Planner step flagged as output_type: image, plus style/brand requirements from Clarification.

Output: One or more generated image files plus a short text description.

V1 Implementation: Single call to an external image-generation API. Output passes through Reviewer for relevance/appropriateness check.

Model Tier: External image-generation provider (e.g. DALL-E or Stability AI).

5.11 Visual Asset Clarification & Approval Flow

Role: Images cannot be checked against an objective checklist. Reviewer, Critic, and Judge are not equipped to judge taste on the user's behalf, so visual assets follow a different path: (1) Clarification Agent asks one style question before Design Agent runs. (2) Generated image returns to the user with three options: approve, reject and regenerate, or revise. (3) Regeneration is capped at 3 attempts per visual asset in V1.

Input: N/A — this is a flow, not a single agent.

Output: User approval or revision instruction.

V1 Implementation: Approval gate on every generated image before project can be marked complete.

Model Tier: N/A.

5.12 Strategy Agent

Role: Think about the why and how to succeed behind a user's goal — business, product, and market thinking that determines whether the thing being built will actually work. Runs after Clarification, before Planner, on project-type requests only.

Input: Clarified requirements + domain hint + Signature Question response + Constraint findings.

Output: Structured strategy brief: goal clarity, MVP scope, key risks, success criteria.

V1 Implementation: Single LLM call with a structured prompt explicitly asking for business/product thinking. Output is a short brief, not a long document.

Model Tier: Mid-tier reasoning.

5.2.1 Goal Discovery Mode

When a user's request is too vague to classify or plan, Clarification Agent enters Goal Discovery Mode. Five layers of discovery, one question per layer at a time:

Layer	Purpose	Example Question
1. Goal Discovery	Understand what the user actually wants to achieve	"Why do you want to build this? What is the real goal?"
2. Problem Discovery	Identify the real problem being solved	"What specific problem will this solve?"
3. Audience Discovery	Clarify who this is for	"Who will use this? What do they currently use instead?"
4. MVP Discovery	Find the minimum viable starting point	"If you could only launch with 3 features, which 3?"
5. Success Discovery	Define what success looks like	"How will you know this worked 90 days after launch?"

5.2.2 Smart Clarification Layer

- **Emotional Context Detection:** reads tone, not just words. Softens language and leads with empathy when stress or urgency is detected.
- **Contradiction Detection:** catches conflicts between later requests and earlier stated constraints. Never silently resolves contradictions.
- **Assumption Surfacing:** identifies unstated assumptions and confirms them before Planner locks in scope.
- **Progressive Disclosure:** questions adapt based on previous answers — irrelevant follow-ups are skipped automatically.
- **Intent Confidence Threshold:** 90%+ → proceed; 70–90% → one confirmation question; below 70% → Goal Discovery Mode.

5.2.3 3-Level Deep Clarification

Level	Focus	What the Agent Does
Level 1 — Surface Goal	What did the user literally say?	Captures the stated goal. Confirms understanding before probing deeper.
Level 2 — Hidden Goal	Why do they want this?	Asks why the goal matters. Surfaces the real outcome the user is trying to achieve.
Level 3 — Reality Check	Are their assumptions realistic?	Identifies gaps between stated goal and reality. Presents as a question, never a statement.

5.2.5 Signature Question Protocol

Every project-type request — without exception — receives exactly one Signature Question before Builder is ever called: **"Why do you believe this is the right solution?"** This single question stops more wrong projects than any amount of technical review. Hard rule: Clarification Agent never blocks a user who insists on proceeding. If the user says 'just build it,' Thinker Core logs the unanswered question, notes the risk in the Strategy Brief, and proceeds. The user's autonomy is always respected.

6. Model & Cost Strategy ★ REVISED

6.1 Model Tiers ★ REVISED

Tier	Used By	Purpose	Pool Size	Provider Diversity
Fast/cheap	Intent, Clarification, Research	High-frequency, low-complexity	6 models	Min 2 providers
Mid-tier reasoning	Planner, Strategy, Reviewer	Structured reasoning	3 models	Min 2 providers
Strongest available	Builder, Judge	Highest-stakes generation	3 models	Min 3 providers
Image generation	Design Agent	Visual asset generation	2 models	Min 2 providers
Multi-model	Consensus Agent	Tie-breaking votes	3 models	Exactly 3 providers

Provider Diversity Rule: within any single agent's model pool, no two models may share the same provider. A provider-wide outage removes at most one model from any pool.

6.2 Resource Protection Law in Practice

- Every pipeline run is logged with token cost per agent.
- The Clarification gate prevents Builder, the most expensive agent, from running against an unclear request.
- Consensus Agent is invoked only when Judge's confidence is borderline — estimated at under 20% of all requests.
- The Reviewer retry loop is capped at one retry to prevent runaway cost on a stuck request.
- Think Credits (Section 21) make cost visible and user-controlled — no silent overruns.

6.3 Direct Chat Path

Not every message needs the full pipeline. If Intent Agent classifies a message as ordinary chat (`next_action: direct_answer`), Thinker Core skips straight to a single direct-answer call and returns. The full pipeline is reserved for project requests.

6.4 When Consensus Actually Fires

Judge Score	Action
Clearly passes all five criteria	Approve directly — no Consensus call
Clearly fails (safety or accuracy critically low)	Reject directly — send back to Builder
Borderline on any criterion	Invoke Consensus Agent for that criterion only
High-risk category (auth, payment, data deletion, permissions)	Consensus fires automatically regardless of score

6.5 Model Identity Concealment

No provider or model name (Claude, GPT, Gemini, or any other) appears anywhere the user can see — not in chat responses, not in Consensus votes, not in error messages, and not in any debug/admin view shipped in V1. Internally, each model slot is referred to by a generic label (e.g. 'Agent A', 'Agent B'). Failure messages describe failures generically ('one of the verification steps didn't respond') — never naming the provider.

6.6 Model Pool Size ★ REVISED

- V1 pool sizes: minimum 3 models per tier for reasoning/generation, minimum 2 for image generation.
- No magic numbers in code — pool size is driven by configuration, not hard-coded constants.
- Per-agent pools, not global: Builder's pool and Reviewer's pool are separate even if they share the same tier.
- Pool health monitoring: a model that fails more than 20% of calls in a rolling 1-hour window is temporarily removed from its pool.
- Beyond V1: pool size becomes fully configurable per agent without a pipeline redesign.

7. Data Flow & State ★ REVISED

7.1 Pipeline State Object ★ REVISED

Each project in progress is tracked as a single state record. All original fields are retained. The following are added to support dynamic routing and resilience:

New Field	Type	Purpose
current_agent	string	Which agent is currently executing
routing_history	array	{agent, next_action, reason, model_used, timestamp} — full trace of every routing decision
loop_counts	object	{planner: int, builder: int, reviewer: int, ...} — per-agent retry counter
failover_log	array	{agent, failed_model, replacement_model, timestamp, peer_audit_result}
peer_audit_flag	boolean	Set to true when a failover model must audit previous output before continuing
resume_from_agent	string null	Set on halt — identifies which agent to restart from on resume
signature_question_response	string null	User's answer to the Signature Question, or null if declined
constraint_findings	object	Output of Constraint Discovery — time, budget, skill, network, existing assets
assumption_flags	array	List of surfaced assumptions and whether each was confirmed or flagged as uncertain

7.2 Memory System

Memory Type	Scope	V1 Storage	Retention Rule
Short-term	Single conversation turn	In-memory / session	Cleared on session end
Project memory	One project across its lifecycle	Database, keyed by project id	Persists until user deletes the project
Long-term preference	User-level, across projects	Database, keyed by user id	Persists indefinitely unless user clears it
Decision Memory	User-level, explicit saved decisions	Database, keyed by user id	Persists indefinitely unless user revokes a decision

Hard rule: passwords, PINs, API keys, and other credentials are never written to any memory layer, even temporarily.

7.3 Decision Memory — Detail

Decision Memory is set explicitly — the user makes a deliberate statement like 'always use React for my projects' and expects the system to follow that rule every time without asking again. When a user states a rule explicitly, Thinker Core writes it to Decision Memory and confirms to the user. On every subsequent project, relevant decisions are injected into Planner Agent's context automatically. If a new request conflicts with a saved decision, Thinker Core surfaces the conflict explicitly rather than silently overriding.

8. Agent Fallibility & Error Containment ★ REVISED

8.1 Why This Section Exists

Every agent is an LLM call, and every LLM call can be wrong. Three distinct failure modes: (1) a single agent gives a wrong verdict, (2) a wrong verdict propagates silently to downstream agents, (3) correlated failure across agents that share the same model family and therefore the same blind spot.

8.2 Containment Rules

Rule	What It Prevents
No single agent's approval is final — Judge approval is always required	A single fallible verdict reaching the user unchecked.
If Reviewer and Critic disagree, Thinker Core routes the disputed point to Consensus rather than silently resolved by whichever agent runs next	Disagreements that it has silently resolved by having the wrong agent run next
Critic's prompt withholds Reviewer's verdict until after Critic has formed its own opinion	Critic anchoring on 'Reviewer already said it's fine' and skipping its own opinion
For high-risk categories (auth, payment, data deletion, permissions), Consensus files a mandatory second review before approving dangerous actions	A shared file is automatically generated by the system, regardless of scale
Every rejection or override is logged with both verdicts and reasoning.	Loss of visibility into how often agents actually disagree.

8.3 What This Does Not Solve

This containment strategy reduces the chance of a bad deliverable reaching the user undetected — it does not guarantee correctness. Consensus Agent's majority vote can still be wrong if a mistake is common enough across providers. No automated review pipeline replaces a human inspecting high-stakes output, particularly around authentication, payments, or anything touching real user data.

8.4 Peer Recovery Protocol ★ NEW

When a model fails mid-task and a replacement model takes over, it follows the Peer Recovery Protocol — it does not simply restart from scratch or blindly continue:

- Step 1 — Failure Detection:** Thinker Core detects failure when a model returns an error, times out after 5 seconds, or returns malformed output. Detection is automatic.
- Step 2 — Context Handoff:** Thinker Core passes the full Pipeline State to the replacement model, including the original agent input, any partial output, and peer_audit_flag set to true.
- Step 3 — Peer Audit:** The replacement model first reviews what the failed model produced. It identifies what is correct, what is incorrect, and what is missing. Audit result is written to failover_log.
- Step 4 — Targeted Continuation:** The replacement model generates only what is needed — keeps correct partial output, fixes incorrect output, completes missing output. Never regenerates correct work.
- Step 5 — Normal Pipeline Resumes:** The replacement model's output is treated identically to a first-attempt output and proceeds through Reviewer, Critic, and Judge as normal.

8.5 Loop Detection Rules ★ NEW

Loop Type	Threshold	Thinker Core Response
Single agent retry (same step)	3 attempts	Halt pipeline. Save state. Notify user.
Planner → Builder → Reviewer full cycle	2 full cycles	Return to Clarification with failure context. Ask user to re-scope.
Clarification loop (vague answers)	4 rounds	Surface the specific unknown directly.
Judge → Consensus → reject cycle	2 cycles	Halt and present best available output with explicit quality caveats.

9. Pipeline Observability & Post-Delivery Feedback

9.1 The Silent Failure Problem

When a user says 'this output is bad,' the problem could have originated at any point in the pipeline. Because agents run sequentially and each feeds the next, a mistake at step 1 propagates silently through all subsequent steps. Without observability, even the founder cannot debug what went wrong.

9.2 Per-Agent Logging (Founder-Visible)

Field	What It Records
agent_name	Which agent ran
input_summary	Short description of what this agent received as input
output_summary	Short description of what this agent produced
duration_ms	How long the agent took in milliseconds
status	success, failed, or skipped
error_detail	If status is failed — error message using generic terms, never model/provider names
confidence	Where applicable — Judge's score, Consensus vote breakdown, Reviewer pass/fail
retry_count	How many times this agent was retried before succeeding or giving up
failover_cost	Token spend from failover calls (separate from normal pipeline cost)
clarification_depth	Which Clarification layers ran (Clarification Agent only)
signature_q_answered	Whether the user answered the Signature Question (Clarification Agent only)

9.4 Post-Delivery Feedback Loop

User's Problem	System Response
Bug in the output	Ask user to describe the bug, send back to Builder with bug description + original plan. Reviewer re-r
Missing feature	Treat as a new small requirement. Run Clarification → Planner (for just the missing piece) → Builder
Style/design not to taste	Follow Visual Asset revision flow (Section 5.11) for images; ask for specific style feedback for code/la
Vague ('it's bad', no detail)	Ask one clarifying question: 'Is there a specific feature that's broken, something missing, or does it lo

9.5 Fix Limits — Cost-Based, Not Count-Based

Fix Type	What Runs	Limit
Small targeted fix (change a color, update a label)	Builder only	Unlimited
Medium fix (fix a bug, add a missing feature)	Builder + Reviewer + Critic	Up to 10 per project session

Fix Type	What Runs	Limit
Full rebuild (start over from requirements)	Full pipeline from Planner onwards	Up to 3 per project session
Rollback (restore a previous version)	Database lookup only — no agent runs	No limit — always free

9.6 Version History

Every time Builder produces an output, Thinker Core saves that version to the database before overwriting the current state. Only the last 10 versions are kept per project in V1. Rollback restores a previous version instantly with no rebuild — rollbacks never count against any fix limit.

9.7 Selective Rollback & Merge

A full rollback restores everything. A selective rollback restores a previous version as the base while keeping specific elements from the current version. Five rules: (1) Full Rollback = database lookup only. (2) Selective Rollback = Clarification clarifies what to keep → previous version restored → Builder applies only selected elements → Reviewer checks merged result. (3) Rollback never destroys history. (4) Every selective rollback creates a new version entry. (5) User-selected changes always override rollback state — Thinker Core never silently resolves a merge conflict.

10. Technology Stack (V1 Recommendation)

Layer	Choice	Reasoning
Frontend	React + Vite	Matches existing developer experience
Backend	Node.js (Express) or Firebase Cloud Functions	Lightweight, matches existing Firebase familiarity
Database	Firestore	Document model fits the pipeline state object well
Auth	Firebase Auth	Reuse existing pattern
LLM Access	Direct API calls to each provider, made only from the Backend	Protect the API keys, matches Resource Protection Law
Payments	Stripe	Webhook-based plan tier updates; Thinker AI never touches raw card data
Hosting	Firebase Hosting or Vercel	Simple, matches existing deployment habits

11. Build Order (Implementation Roadmap) ★ REVISED

Phase	Deliverable	Status
Phase 1	Thinker Core skeleton + direct chat path (passthrough)	Unchanged
Phase 1.5 ★ NEW	Dynamic Routing Engine: routing table, next_action handling, loop_counts, loop_detection, loop_protection, loop_recovery, loop_protection (Section 5.2.1)	Unchanged
Phase 1.6 ★ NEW	Model Pool Manager — failover detection (5s timeout), pool health monitoring, New-Router-Complete (Section 5.2.2)	Unchanged
Phase 2	Intent Agent + Clarification Agent (basic) + hard gate	Unchanged
Phase 2.5	Smart Clarification Layer (Section 5.2.2)	Unchanged
Phase 2.6	3-Level Deep Clarification + Signature Question (Sections 5.2.3, 5.2.5)	Unchanged
Phase 2.7	Constraint Discovery + Assumption Discovery (Section 5.2.4)	Unchanged
Phase 3	Planner Agent + Builder Agent (single-step)	Unchanged
Phase 4	Reviewer Agent + one-retry loop back to Builder	Unchanged
Phase 5	Critic Agent + Judge Agent (Judge decides alone)	Unchanged
Phase 6	Consensus Agent + borderline-routing from Judge	Unchanged
Phase 7	Research Agent (optional step augmentation)	Unchanged
Phase 8	Design Agent + Builder/Design routing (Section 4.3)	Unchanged
Phase 8.5	Strategy Agent (Section 5.12)	Unchanged
Phase 9	Project memory + long-term preference + Decision Memory	Unchanged
Phase 10	Domain picker wiring (Section 4.4)	Unchanged
Phase 10.5	Per-agent logging (Section 9.2) + post-delivery feedback (Section 9.4)	Unchanged
Phase 10.6 ★ NEW	Resilience observability: failover_log dashboard, pool health alerts, routing_history trace viewer	Unchanged
Phase 10.7 ★ NEW	Think Credits integration: credit deduction engine, balance checks, confirmation prompts (Section 21)	Unchanged
Phase 11	Polish: UI, error states, retry UX, and the launch checklist (Section 12)	Unchanged

12. V1 Launch Checklist

- Stable chat experience: the direct path works without errors
- Full agent pipeline completes end-to-end on at least three different project types (app, website, game)
- Strategy Agent produces a meaningful strategic brief, visibly reflected in Planner's output
- Decision Memory saves an explicit user rule, applies it automatically, and surfaces conflicts correctly
- Design Agent produces a usable image; approve / reject / revise all work correctly; 3-attempt cap enforced
- Clarification gate cannot be bypassed
- Smart Clarification correctly detects emotional tone and adjusts question framing
- Contradiction Detection catches at least one case where user's later message contradicts an earlier constraint
- Signature Question fires on every project-type request; response stored in Pipeline State
- Intent Confidence Threshold correctly routes: high → proceed, mid → confirm, low → Goal Discovery Mode
- No credentials ever appear in stored memory (manual audit)
- Consensus Agent tested with intentionally borderline cases
- High-risk categories tested to confirm Consensus fires automatically
- Token cost per request is logged and reviewed
- Final response is always a single merged answer, never raw agent dumps
- 5-stage output pipeline works end-to-end (Section 16)
- Live Preview renders correctly for web projects
- Language detection works: user writes in any Tier 1 language and all agent responses are in that language
- RTL language (Arabic or Hebrew) renders correctly
- Per-agent log entries written for every pipeline run; failed run traceable to specific agent
- Cost-based fix limits enforced: small unlimited, medium capped at 10, full rebuild capped at 3
- Version history saves correctly; rollback restores without Builder call; 10-version limit enforced
- Selective rollback works correctly
- Domain picker measurably changes pipeline output
- No provider/model names reachable from any user-facing surface including error states
- Dynamic routing engine tested: all next_action values handled correctly
- Failover tested: model timeout triggers pool reassignment within 5 seconds
- Peer Recovery Protocol tested: replacement model audits and continues partial work correctly
- Loop detection tested: exceeding thresholds triggers halt and state save
- Think Credits deducted correctly per action; no silent overruns; confirmation shown above 3 credits

- Stripe webhook correctly updates plan tier on payment and downgrade
- Free Trial 50-credit one-time limit enforced; no monthly refill

13. Explicitly Out of Scope for V1

- Parallel multi-step Builder execution
- Automatic re-planning when a step fails repeatedly (Peer Recovery handles model-level failure; plan-level re-planning is V2)
- Native mobile app (Play Store launch can wrap the web app initially)
- Organisational/team-shared memory
- Autonomous multi-day project system with no human checkpoint
- Native Thinker foundation model (V1 uses third-party model APIs only)
- 3D model generation
- Animation and video generation
- UI/UX design as a distinct discipline — Design Agent produces static images only, not Figma-format output

14. Open Decisions for Founder Review

1. Consensus uses exactly three providers — confirmed for V1, configurable beyond three in later versions.
2. Reviewer retry is capped at one loop — confirm this is acceptable or should be configurable.
3. Research Agent is optional and skippable per step — confirm this matches intended behaviour.
4. Technology stack assumes continuity with existing Firebase-based project — confirm versus fresh stack.
5. High-risk category list (auth, payments, data deletion, permissions) — confirm this list is complete.
6. Design Agent's image provider not yet chosen between DALL-E or Stability AI — confirm based on cost, quality, and licensing.
7. Free Trial credit amount set at 50 (one-time) — confirm this is the right conversion balance.
8. Pro price set at \$19/month and Founder at \$59/month — confirm based on competitive positioning.

15. Goal Intelligence System

This section brings together the highest-level thinking layer of Thinker AI — the system that determines whether a user's goal is valid, well-formed, and worth building, before any technical work begins. This is what makes Thinker AI a thinker, not just a builder.

15.1 Goal Discovery

Goal Discovery is the process of understanding what the user actually wants to achieve. When a user's goal is unclear, Goal Discovery Mode activates (Section 5.2.1) and works through five layers of questions — one at a time, conversationally. Goal Discovery is not an obstacle — it is the product. A user who spends 5 minutes in Goal Discovery before building is far more likely to end up with something useful.

15.2 Problem Discovery

Problem Discovery is focused specifically on identifying the real problem being solved. If the problem is not real, or not specific enough, no amount of good engineering will produce a successful product. Thinker AI asks: 'What specific problem will this project solve?' and follows up until the answer is specific.

15.3 Idea Validation

After Goal and Problem Discovery, Strategy Agent runs an Idea Validation covering four questions: (1) Is the problem real? (2) Is this idea the right solution? (3) Is this the right time? (4) Why this team? The output is not a go/no-go decision — it is a brief that sharpens the plan Planner receives.

15.5 Founder Mode

Activated when a user is thinking through an entire product, business, or startup idea. Can be triggered explicitly ('Activate Founder Mode') or detected automatically when the request involves market, revenue, competition, or business model language. In Founder Mode, Strategy Agent expands its scope to cover: Problem → Market → Competition → Business model, full risk landscape, and a Go/Caution/No-Go assessment (always advisory, never a block).

16. Output Pipeline — Preview Before Build

Instead of building everything silently and delivering a ZIP at the end, Thinker AI follows a 5-stage pipeline where the user sees and approves each stage before the next begins. This implements the Preview Before Build Law (Section 3).

Stage	Name	What User Sees	User Action
1	Thinking Summary	Goal understanding, strategy brief, proposed approach	Confirm or correct before anything is built
2	Blueprint Preview	Full plan: screens/components, tech stack, estimated app complexity	Approve, remove/add items, or request changes
3	Build Preview (live)	Real-time progress + live rendered UI updating as each screen is completed	Approve first completed mandatory, or change directly
4	Approval	Full app visible in preview panel	Approve, change something specific, or go back to blueprint
5	Final Output	Project Summary, Live Preview Link, ZIP download, Archived and Notes	

Simple requests skip stages 2 and 4. The user is never surprised by what was built — even for simple outputs, Thinker Core shows a one-line summary and waits for brief confirmation before running Builder.

17. Multi-Language Support

17.1 Core Principle

Thinker AI detects the user's language automatically from their first message and responds in that same language throughout the entire session — all agent responses, clarification questions, strategy briefs, progress indicators, error messages, and final summaries. The user never has to specify their language or change a setting.

17.2 What Gets Translated vs. What Stays in English

Content Type	Language Used	Reason
All conversation with user	User's language	This is the product interface
Clarification questions, strategy brief, User's Display	User's language	User must understand these clearly
Progress indicators, error messages	User's language	UX must be native-feeling
Generated code (variables, functions, classes)	English	Industry standard; must be readable by any developer
Code comments	User's language (if requested) or English (default)	English (default) via Decision Memory
README and documentation	User's language	User needs to understand their own project
Folder/file names, API endpoint names	English	Compatibility and portability

17.4 Supported Language Tiers

Tier	Languages	Quality in V1
Tier 1 — Full support	English, French, Spanish, German, Portuguese, Italian, Dutch, Japanese, Korean, Chinese (Simplified), Arabic, Hebrew, Urdu, Farsi, and others	High quality
Tier 2 — Good support	50+ additional languages including Urdu, Malay, Thai, Swahili, Czech, Slovak, Norwegian, Danish, Polish, Vietnamese, Indonesian, and others	Good quality
Tier 3 — Basic support	Remaining languages supported by the underlying model provider	Response quality depends on the model's training

RTL Languages (Arabic, Hebrew, Urdu, Farsi, and others): fully supported in all user-facing text. The UI must render these correctly — right-aligned text, mirrored layout where appropriate. V1 must be tested with at least one RTL language before launch.

18. Paid Plan Architecture ★ NEW

All plan enforcement happens inside Thinker Core, immediately after the incoming request is received and before any agent is invoked. No billing or quota logic exists inside individual agents.

18.1 Plan Tiers Overview

Tier	Price	Think Credits	Target User
Free Trial	\$0	50 (one-time, never refills)	Individuals exploring the product
Pro	\$19/month or \$199/year	1,500/month	Freelancers, indie builders
Founder	\$59/month or \$589/year	5,000/month	Power users, startup founders
Enterprise	Custom	Custom	Organisations 10+ seats (V2)

18.2 Feature Gate Matrix

Feature	Free Trial	Pro	Founder
Direct Chat	✓	✓	✓
Intent Agent	✓	✓	✓
Basic Clarification	✓	✓	✓
Smart Clarification (3-Level Deep, Signature Q)	—	✓	✓
Strategy Agent	—	✓	✓
Founder Mode	—	✓	✓
Planner Agent	✓	✓	✓
Research Agent	—	Limited	Unlimited
Builder Agent	Credit-limited	Credit-limited	Credit-limited
Reviewer + Critic + Judge	Reviewer + Judge only	✓	✓
Consensus Agent	—	—	✓
Design Agent	—	✓	✓
5-Stage Output Pipeline	Stage 1+5 only	Full	Full
Version History & Rollback	3 versions	10 versions	25 versions
Decision Memory	—	—	✓
Thinking Level: High (■)	—	✓	✓
Thinking Level: Consensus (■)	—	—	✓
Priority Queue	—	—	✓

Feature	Free Trial	Pro	Founder
Extra credit rate	\$0.10/credit	\$0.07/credit	\$0.05/credit

18.3 Enforcement Architecture

Thinker Core enforces plan rules in 5 steps before any agent runs:

Step 1 — Authenticate: Read user's plan tier from session token. Plan tier is refreshed on billing events, never re-fetched mid-pipeline.

Step 2 — Feature gate check: Check whether the requested pipeline path is permitted for this tier. A gated feature returns a structured error with upgrade prompt — the pipeline does not partially execute.

Step 3 — Credit balance check: Read Think Credit balance. If insufficient for estimated session cost, block and show top-up or downgrade options.

Step 4 — Concurrent run check: If user already has the maximum active pipeline runs for their tier, queue the new request rather than rejecting it.

Step 5 — Pipeline proceeds: If all checks pass, pipeline runs normally. Plan tier is injected into Pipeline State for observability logging only.

18.4 Billing Integration

- Stripe Checkout handles all payment collection. Thinker AI never touches raw card data.
- On successful payment, Stripe sends a webhook to a dedicated billing endpoint (not /api/pipeline). The billing endpoint updates the user's plan_tier in Firestore and refreshes the session token.
- Downgrade takes effect at end of current billing period. User retains access until then. On downgrade, Pro-created projects are preserved but gated features become read-only.
- Failed payment: user enters 'grace' state — 7 days of continued access while Stripe retries. After 7 days without resolution, tier reverts to Free Trial automatically.
- Refunds handled entirely in Stripe. The plan_tier change is the only signal Thinker Core needs.

19. Resilience Architecture — Continuous Delivery Guarantee

★ NEW

The Continuous Delivery Guarantee: Thinker AI never holds a user's work hostage to a single model, a single provider, or a temporary outage. When any model fails, another model takes the task, audits what was done, and continues. The user's work always moves forward.

19.1 No Single Point of Failure

- Minimum pool size: 3 models for all reasoning and generation agents.
- Provider diversity: no two models in the same agent's pool may share a provider.
- Pool configuration is static per deployment — Thinker Core takes a pool_health_snapshot at pipeline start.
- If a pool's health snapshot shows fewer than 2 available models at pipeline start, Thinker Core notifies the operator (not the user) and proceeds with reduced redundancy.

19.2 Failure Detection Timeline

Time	Event	Thinker Core Action
0s	Model call initiated	Start 5-second timeout clock
0–5s	Model responds normally	Continue pipeline — no action needed
5s	No response received	Mark model as failed. Assign next model from pool.
5–7s	Replacement model initialises	Peer audit begins (Section 8.4 Steps 2–3)
7s+	Replacement model continues task	Pipeline resumes — user sees no pause
All pool models fail		Loop detection triggers halt. State saved.

19.3 User Communication During Failures

Situation	Message to User
Failover resolves within 7s	No message. Pipeline continues normally.
Failover takes 7–30s	'Still working on this — taking a moment longer than usual.'
Partial completion, halt triggered	'I've completed [X of Y steps]. I ran into a difficulty on the next step — your work so far is saved. Would you like to try again?'
All models unavailable	'I wasn't able to complete this step right now — the service I need isn't responding. Your work is saved. I'll try again in a moment.'

Hard rule: no provider or model name appears in any user-facing message, including failure messages (Section 6.5).

19.4 Failover Cost Rules

- Failover additional model calls are absorbed by the operator — never charged to the user (0 credits).
- Peer audit scope is minimal — replacement model audits only the failed model's partial output.
- Targeted continuation — replacement generates only what is missing or incorrect; correct partial output is reused.
- Failover token spend is logged under `failover_cost` in the per-agent log (Section 9.2), separate from normal pipeline cost.
- A model failing more than 20% of calls in a rolling 1-hour window is temporarily removed from its pool.

20. Dynamic Routing Engine — Implementation Specification

★ NEW

This section specifies the internal design of the routing layer inside Thinker Core introduced in Section 4.1. It is an implementation guide for the developer, not a user-facing feature description.

20.1 Routing Engine Responsibilities

- Read `next_action` from every agent output before passing control forward.
- Maintain the `routing_history` array in Pipeline State — one entry per routing decision.
- Enforce `loop_count` thresholds (Section 8.5) and trigger halt when exceeded.
- Manage the Model Pool Manager: track pool health, execute failovers, trigger peer audits.
- Never expose routing internals to agents — agents receive only their defined inputs.
- Never expose routing internals to users — users see only the result.

20.2 Routing Decision Logic (pseudocode)

```
function routeAfterAgent(agent_output, pipeline_state):  
    action = agent_output.next_action  
    log_routing_decision(agent_output, pipeline_state)  
  
    if action == 'proceed':  
        return dispatch(next_in_sequence(pipeline_state.current_agent), pipeline_state)  
    if action == 'retry':  
        pipeline_state.loop_counts[current_agent] += 1  
        if loop_counts[current_agent] > 3: return trigger_halt(pipeline_state)  
        return dispatch(current_agent, pipeline_state)  
    if action == 'replan':  
        pipeline_state.loop_counts['planner_cycle'] += 1  
        if loop_counts['planner_cycle'] > 2: return dispatch('clarification', pipeline_state)  
        return dispatch('planner', pipeline_state_with_failure_context)  
    if action == 'clarify': return dispatch('clarification', pipeline_state)  
    if action == 'escalate': return dispatch('consensus', pipeline_state)  
    if action == 'direct_answer': return dispatch('direct_llm_call', pipeline_state)
```

20.3 Model Pool Manager Logic (pseudocode)

```
function callAgentWithFailover(agent_name, input, pipeline_state):  
    pool = get_pool(agent_name) # from pool_health_snapshot  
    for model in pool.available_models:  
        try:
```

```
result = call_model(model, input, timeout=5s)
if is_valid_output(result): return result
except TimeoutError, ModelError:
    log_failover(agent_name, model, pipeline_state)
    next_model = pool.next_available(exclude=failed_models)
    if next_model:
        input_with_audit = add_peer_audit_flag(input, partial_result)
        # recurse with updated exclusion list
    else:
        return trigger_halt(pipeline_state) # pool exhausted
```

20.4 What the Routing Engine Is Not

- Not an LLM call — the routing engine is plain backend logic (conditional branching on next_action values).
- Not visible to agents — agents receive their defined inputs and return their defined outputs.
- Not a replacement for good agent design — it handles failures gracefully, not poor prompts.
- Not a billing layer — plan tier checks (Section 18.3) run before the routing engine, not inside it.

21. Think Credits, Pricing & Thinking Levels ★ NEW

The Thinker AI Golden Rule

More Thinking
= More Questions
= More Validation
= More Credits

Users do not pay for answers. They pay for the depth of thinking behind the answer. This principle aligns every pricing tier, every credit cost, and every agent behaviour in Thinker AI into a single coherent system.

21.1 Why Every Interaction Costs Credits

Consumer AI products like ChatGPT or Claude.ai own their foundation models — their marginal cost per message approaches zero. Thinker AI is architecturally different: every agent call invokes third-party model APIs (Section 6.1). Each call has a real monetary cost paid to the provider. The Think Credit system makes this cost visible, user-controlled, and sustainable — every interaction, including simple chat, costs at least 1 credit.

21.2 Thinking Levels

Level	Icon	Purpose	Agents Active	Base Cost
Low Thinking	■	Fast answers, simple chat	Intent + Direct Answer only	1 credit
Medium Thinking	■	Analysis, comparisons, recommendations	Intention Clarification + Planner + Builder + Reviewer	8 credits
High Thinking	■	Deep planning, full project builds	Full pipeline (all 12 agents)	10 credits
Consensus Thinking	■	Multi-model validation, startup ideas	Full pipeline + Consensus (3 providers)	30 credits

Thinker Core auto-selects the Thinking Level during Intent Agent classification. The user can accept, downgrade, or upgrade before the pipeline runs. Credits are never deducted without user confirmation for sessions costing more than 1 credit.

21.3 Plan Restrictions by Thinking Level

Thinking Level	Free Trial	Pro	Founder
■ Low Thinking	✓	✓	✓
■ Medium Thinking	✓	✓	✓
■ High Thinking	—	✓	✓
■ Consensus Thinking	—	—	✓

21.4 Clarification & Goal Discovery Depth by Thinking Level

Layer	■ Low	■ Medium	■ High	■ Consensus
Clarification style	Direct (0–2 Q)	Smart (3–5 Q)	Deep (5–10 Q)	Deep + Challenge (Dynamic)
Goal Discovery	—	Light	Full	Full
Problem Discovery	—	Light	Full	Full
Audience Discovery	—	Light	Full	Full
MVP Discovery	—	Optional	Full	Full
Success Discovery	—	Optional	Full	Full
Assumption Discovery	—	—	Conditional	Full
Constraint Discovery	—	Optional	Full	Full
Alternative Discovery	—	—	Conditional	Full
Reality Challenge	—	—	Light	Strong
Founder Mode	—	—	—	✓

21.5 Credit Cost Table

Action	Credits	Notes
Simple chat / direct answer	1	Per message. Always applies.
Clarification reply (per round)	2	Each round of questions + user answer = 2 credits
Research Agent (per step)	3	Only when Planner flags a step as research-needed
Medium analysis	4	Medium Thinking sessions
Image generation (per image)	10	Per image, per attempt — including revision attempts
High Thinking session (base)	10	Charged once at session start
Consensus session (base)	30	Charged once at session start
Small fix (Builder only)	1	Color change, label update, minor edit
Medium fix (Builder + Reviewer + Critic)	5	Bug fix, missing feature
Full rebuild (Planner onwards)	50	Complete restart from planning
Rollback (any version)	0	No LLM call — database lookup only
Selective rollback + merge	3	One small Builder call + one Reviewer call
Failover (model substitution)	0	Absorbed by operator — never charged to user

Confirmation rule: for any single action costing more than 3 credits, Thinker Core shows the cost before executing and waits for user confirmation.

21.6 Worked Examples

User Request	Level	Flow	Total Cost
React ■■■?	■ Low	Intent → Direct Answer	1 credit
React vs Vue — ■■■■ Startup ■■■■? ■■■■	■■■ Medium	Intent → Clarification (2 Q) → Planner → Builder → Reviewer	3 + 2 + 4 = 9 credits
Food Delivery Startup ■■■■ ■■■■ ■■■■	■■■ Medium	Intent → Clarification (3 rounds, Deep Discovery) → Strategy → Planner → Builder → Reviewer	15 + 5 + 5 + 5 + 5 = 35 credits
■■■■ Startup Idea validate ■■■■ Consensus	■■■ Medium	Intent → Clarification (5 rounds, all layers) → Strategy → Planner → Builder → Reviewer → Pipeline	90 + 10 + 10 + 25 + 75 = 220 credits

21.7 Credit Enforcement Rules

- **Check before dispatch:** Thinker Core checks credit balance before dispatching any agent. Insufficient balance = pipeline does not start.
- **Confirm above 3 credits:** any single action costing more than 3 credits shows: action name, credit cost, current balance, balance after. User must confirm.
- **Deduct on completion, not on start:** credits deducted when an agent successfully completes, not when dispatched. Failed calls that trigger failover are not charged.
- **Failover is always free:** model substitution events never consume user credits.
- **Partial session crediting:** if pipeline halts mid-run, credits charged only for agents that successfully completed. Thinking Level base cost still applies if session was initialised.
- **Low balance warning:** balance below 20 credits shows a persistent indicator. Balance at 0 shows top-up prompt before any new session.
- **No silent overruns:** if mid-session estimated remaining cost exceeds remaining balance, Thinker Core pauses and asks the user to confirm or stop.

The Thinker AI Golden Rule

More Thinking
= More Questions
= More Validation
= More Credits

Every credit a user spends is a unit of reasoning, validation, and challenge applied to their goal. The more they invest in thinking, the less likely they are to build the wrong thing. This is Thinker AI's value proposition — and its pricing model.

Thinker AI V1 Architecture Spec

Addendum 3 — Sections 22 & 23

File Upload Architecture | UI/UX Specification

June 2026

Section 22 defines how Thinker AI handles every type of file a user can upload — documents, code, images, data files, and ZIP archives — including a new Document Agent (the 13th agent role in V1). Section 23 specifies the complete UI/UX structure: every screen, every settings panel, every profile option, and every in-chat control a user can interact with.

22. File Upload Architecture

Thinker AI accepts any file a user uploads. Thinker Core identifies the file type, routes it to the correct agent, and determines what to do with it — the user does not need to specify 'review this code' or 'analyze this PDF.' Thinker reads the file and asks the right question.

22.1 Core Principle

File Upload Pipeline: File arrives → Thinker Core identifies type → Security scan → Route to correct agent → Agent reads & understands → Clarification Agent asks what to do → Pipeline runs → Deliver result.

22.2 Supported File Types & Routing

Category	Formats	Max Size	Routed To	Credit Cost
Documents	PDF, DOCX, TXT, MD	50 MB	Document Agent (Section 22.7)	3 credits per file
Code files	JS, TS, PY, HTML, CSS, JSON, XML, and all code extensions	10 MB	Review extension → Judge → Builder (if fix requested)	5 credits review 10 credits fix
Images	PNG, JPG, JPEG, SVG, WEBP, GIF	20 MB	Design Agent	2 credits analysis 15 credits edit 20 credits reference-gen
Data files	CSV, XLSX, JSON	25 MB	Research Agent	3 credits per file
Archives	ZIP, TAR	100 MB	Thinker Core (extract & re-route each file)	5 credits extract + per-file costs

Blocked file types: .exe, .bat, .sh, .dll, .bin, .dmg, and any executable format. These are rejected before processing begins — no credits are charged.

22.3 ZIP / Archive Upload Flow

A ZIP upload is treated as a full project hand-off. Thinker Core extracts the archive, maps the structure, and then asks the user what to do — rather than assuming.

- ZIP arrives → Thinker Core extracts to temporary in-memory workspace
- File inventory: count, types, folder structure mapped and shown to user
- Each file routed by type (code → Reviewer, images → Design Agent, docs → Document Agent)
- Project Memory populated with full project context
- Clarification Agent asks: "What would you like to do with this project?"
- Options: Review entire codebase | Add a feature | Fix a bug | Rebuild | Analyse only
- Selected action → normal pipeline runs with full project context
- Output: modified ZIP with same folder structure + changelog

ZIP size limit: 100 MB. If the extracted content exceeds 500 files, Thinker Core summarises the structure and asks the user to select which subfolder or module to focus on first.

22.4 Code File Upload Flow

Step	What Happens
1. Upload	Code file arrives. Thinker Core identifies language from extension and content.
2. Auto-review	Reviewer Agent reads the code. Critic Agent independently looks for issues. Both run automatically — no user prompt needed.
3. Findings shown	User sees: language detected, file size, issues found (by severity), strengths noted.
4. User chooses action	'Escalate issues' 'Fix only critical issues' 'Explain what this code does' 'Add a feature to this code' 'Rewrite this in [language]'
5. Pipeline runs	Builder executes the chosen action. Judge approves before result is shown.
6. Output	Modified code file + diff showing what changed + brief explanation of changes.

Security rule: if code contains patterns matching known malware, cryptomining, data exfiltration, or shell injection, Thinker Core warns the user explicitly and does not build or improve the code. The warning names the pattern category but does not provide detail that would help weaponise it.

22.5 Image Upload & Editing Flow

Images can be uploaded for three distinct purposes. Thinker Core identifies the intent from the user's message alongside the upload:

Operation	User Intent Signal	Design Agent Action	Credits
Analyse only	'What is in this image?' / no edit instructions	Vision model describes content, objects, colours, style	2
Edit / remove object	'Remove the background' / 'Delete the image on the left' / 'Make the sky blue' / 'Delete object'	Image inpainting, object masking, API applied	15
Style transfer	'Make this cartoon style' / 'Convert to sketch' / 'Simplify image'	Style transfer, GANs applied to full image	15
Reference-based generation	'Use this logo's style to make a banner' / 'Extract colour palette'	Style extraction, colour palette, generate image. New asset generated using extracted	20
Upscale / enhance	'Make this higher resolution' / 'Sharpen image'	Upscaling API applied	8

All image edits follow the Visual Asset Approval Flow (Section 5.11): user sees result and can approve, revise, or regenerate. Regeneration capped at 3 attempts per edit operation.

22.6 Document Upload Flow (PDF, DOCX, TXT)

User Intent	What Happens	Output
Summarise	Document Agent extracts key points, sections, conclusions	Structured summary in user's language
Analyse / critique	Document Agent reads for logic, gaps, contradictions	Strengths/weaknesses, specific recommendations
Extract data	Document Agent identifies tables, numbers, lists, names	Structured data export (JSON or table)
Answer questions about document	Document added to context. User asks questions.	Direct answers with page/section references
Use as project input	Document becomes Clarification input — Thinker builds alongside user with document as requirements source	
Translate	Content extracted, translated to user's target language	Translated document in same format

22.7 Document Agent — 13th Agent Role ★ NEW

Document Agent is a new agent role added to handle file-based input. It is the 13th agent in V1, joining the original twelve.

Field	Specification
Role	Read, parse, and extract structured content from any document file. Convert unstructured document content into structured content.
Input	Uploaded file (PDF, DOCX, TXT, MD, CSV, XLSX) + user intent + conversation context
Output	Structured content object: {summary, key_points, sections[], data_tables[], entities[], language_detected, page_count}
V1 Implementation	Two-step: (1) file parsing library extracts raw text and structure; (2) single LLM call converts raw extraction into structured content.
Model Tier	Fast/cheap — document parsing is mostly rule-based; LLM call is for structuring, not reasoning.
next_action output	Same as all agents: proceed clarify retry halt
Does NOT do	Does not generate new content, write code, or make decisions. It reads and structures — other agents act on it.

22.8 File Security Rules

Executable files blocked: Files with extensions .exe, .bat, .sh, .dll, .bin, .dmg, .apk, .ipa are rejected at upload. No credits charged. User told why.

No file stored permanently: Uploaded files live in temporary in-memory workspace for the duration of the session. They are deleted when the session ends or the pipeline completes. They are never written to persistent storage.

File content stays in context only: File content is sent to LLM providers as context in API calls — the same way a chat message is. It is not uploaded to a separate file storage service. The provider's data handling rules apply (same as any message content).

Malicious code detection: Code files are scanned for patterns matching: shell injection, data exfiltration, cryptomining, known malware signatures. Detection triggers a warning and blocks the build action. The scan runs before any LLM call.

File size enforced pre-pipeline: Size limits are checked before the file enters the pipeline. An oversized file is rejected immediately with a clear message — no credits are charged.

No credentials in files: If a code file contains what appears to be an API key, password, or token (detected by pattern matching), Thinker Core warns the user before proceeding and redacts the value from LLM context. Hard rule: credentials are never sent to model providers even as part of a code review.

22.9 File Upload Credit Cost Summary

Operation	Credits	When Charged
Document analysis (per file)	3	When Document Agent completes successfully
Code review (per file)	5	When Reviewer + Critic complete
Code fix / improvement	10	When Builder completes the fix
ZIP extract + structure map	5	When Thinker Core maps the structure
Image analysis (vision)	2	When Design Agent describes the image

Operation	Credits	When Charged
Image edit / object removal	15	When Design Agent returns edited image
Image style transfer	15	Per attempt
Reference-based image generation	20	Per attempt
Image upscale / enhance	8	When upscaled image is returned
Data file (CSV/XLSX) analysis	3	When Research Agent structures the data
Full project from ZIP	5 (extract) + no multiple charges	Over pipeline
Blocked file (executable, oversized)	0	Never charged — rejected pre-pipeline

23. UI/UX Specification

This section defines every screen, panel, and interactive element in Thinker AI's user interface. It covers the main chat view, navigation, profile, settings, and all in-conversation controls. Design implementation follows the frontend-design skill conventions; this section defines what exists, not how it looks.

23.1 Overall App Layout

Zone	Location	Contents
Left Sidebar	Fixed, collapsible	Logo, Domain Picker, Conversation History, New Chat button, Upgrade CTA (Free/Pro u
Main Chat Area	Centre, scrollable	Message thread, Thinking Level indicator, progress stages, output panels
Input Bar	Bottom, fixed	Text input, file upload button, Thinking Level selector, send button, credit balance chip
Right Panel	Slide-in on demand	Project details, version history, pipeline log (founder mode), settings shortcuts
Top Bar	Fixed, minimal	Current conversation title (editable), domain label, user avatar, notifications bell

23.2 Left Sidebar — Detail

Domain Picker	Conversation History
• CEO Agent (default — general judgment)	• Today / Yesterday / This Week / Older groups
• Coding	• Search conversations (full-text)
• DevOps	• Pin a conversation to top
• Security	• Rename conversation
• Music	• Delete conversation
• Design	• Export conversation (PDF or TXT)
• Data / Analytics	• Star / favourite a conversation
• Marketing	• Project badge (if conversation has a build)
• Finance	• Last message preview
• Legal	• Credit spend shown per conversation
• Research	
• Education	
• Healthcare	
• + Custom domain (Pro/Founder only)	

23.3 Input Bar — Detail

Input Bar Controls

- Text input field — multiline, auto-expands, supports paste of code/URLs
- File upload button (paperclip icon) — opens file picker, supports drag-and-drop
- Thinking Level selector — tap to cycle , or long-press for picker
- Send button — disabled until input is non-empty or file is attached
- Credit cost preview chip — shows estimated credits before sending (updates as Thinking Level changes)
- Voice input button (V2 — not in V1 scope)
- Emoji / reaction button (V2)

23.4 Chat Message Types

Message Type	Visual Treatment	Contents
User message	Right-aligned, user colour	Text, attached file name/thumbnail
Thinker direct answer	Left-aligned, clean	Text response, formatted markdown, code blocks
Thinking Level indicator	Inline badge before response	 icon + level name + credit cost
Stage 1 — Thinking Summary	Expandable card	Goal understood, strategy brief, proposed approach, Confirm/Edit buttons
Stage 2 — Blueprint	Expandable card with checklist	Step list with checkboxes, tech stack, Approve/Edit buttons
Stage 3 — Build Progress	Split view: progress left, preview right	Live weight ticker + rendered UI preview, first-screen approval prompt
Stage 4 — Approval	Full-width card with 3 CTAs	Looks good / Change something / Go back to blueprint
Stage 5 — Final Output	Delivery card	Summary, Live Preview link, ZIP download, Architecture Notes
Clarification questions	Distinct question card	Numbered questions, input fields inline, Submit button
Credit confirmation	Modal overlay	Action name, credit cost, balance before/after, Confirm/Cancel
Upgrade prompt	Inline banner	What was gated, which plan unlocks it, Upgrade CTA
Error / halt message	Amber card	What happened, what is saved, options to continue
File received card	Left-aligned, file icon	File name, type, size, detected language/content type, available actions
Image output card	Left-aligned, image preview	Generated/edited image, Approve/Revise/Regenerate buttons, attempt counter
Version history card	Slide-in panel	Version list with timestamps, Restore button per version, diff preview

23.5 Profile Page

Accessed via the user avatar in the top bar. Contains all personal account information.

Profile — Identity	Profile — Activity
• Display name (editable)	• Total projects built (count)
• Email address (editable, requires re-verification)	• Total Think Credits spent (all time)
• Profile photo (upload, crop, remove)	• Favourite domains used
• Username / handle (unique, used in shared links)	• Most used Thinking Level
• Bio / description (optional, 160 chars)	• Total conversations
• Preferred language (overrides auto-detect)	• Total files uploaded
• Date joined	• Projects shared publicly
• Account type badge (Free / Pro / Founder)	• Member since date

Profile — Decision Memory	Profile — Linked Accounts
• View all saved decisions (list)	• Google (for sign-in)
• Edit a saved decision	• GitHub (for code import — V2)
• Delete / revoke a decision	• Figma (for design import — V2)
• Add a decision manually	• Disconnect a linked account
• Decision applies to: all projects / specific domain	• Add new linked account
• Last applied: date shown per decision	

23.6 Settings Page

Accessed via Settings in the left sidebar or top bar menu. Organised into eight tabs.

Settings Tab: General
• App language (overrides auto-detect for UI elements)
• Response language (same as user language / always English / ask each time)
• Default Thinking Level (auto / always Low / always Medium / always High)
• Default domain (none / pick one)
• Code comment language (English / user language / ask each time)
• Date & time format
• Timezone

- App theme (Light / Dark / System)

- Font size (Small / Medium / Large)

Settings Tab: Thinking & Pipeline

- Auto-select Thinking Level (on/off)

- Show credit cost before sending (on/off — default on)

- Show pipeline progress stages (on/off)

- Show Thinking Level badge on every response (on/off)

- Signature Question (on/off — default on, can be disabled on Pro/Founder)

- Goal Discovery depth (auto / always full / always light)

- Strategy Agent (on/off per domain)

- Founder Mode auto-detect (on/off)

- Multi-language in pipeline (on/off — default on)

Settings Tab: Memory

- Project Memory (on/off)

- Long-term Preference Memory (on/off)

- Decision Memory (on/off)

- Clear all Project Memory

- Clear all Long-term Memory

- Clear all Decision Memory

- Export all memory (JSON)

- Memory retention period (30 / 90 / 180 / forever)

Settings Tab: Files & Uploads

- Default action on code upload (auto-review / ask each time)

- Default action on document upload (summarise / ask each time)

- Default action on image upload (analyse / ask each time)

- Default action on ZIP upload (map & ask / full review)

- Auto-scan code for security issues (on/off — default on)
- Warn before sending file content to LLM providers (on/off)
- Max file size limit (use platform default / set custom lower limit)

Settings Tab: Credits & Billing

- Current plan name + renewal date
- Think Credits balance (remaining this month)
- Credits used this month (with breakdown by action type)
- Buy extra credits button
- Upgrade / downgrade plan
- View billing history (invoice list)
- Download invoice (per month)
- Cancel subscription
- Low balance alert threshold (set custom — default 20 credits)
- Auto top-up (on/off — buy X credits when balance drops below Y)
- Payment method (view / update)

Settings Tab: Notifications

- Pipeline complete (push / email / none)
- Build failed / halted (push / email / none)
- Low credit balance alert (push / email / none)
- New feature announcements (email / none)
- Weekly usage summary (email / none)
- Billing / payment notifications (email — always on, cannot disable)
- Upgrade reminders (on/off)

Settings Tab: Privacy & Security

- Change password
- Two-factor authentication (on/off, TOTP or SMS)

- Active sessions (list, revoke any session)
- Data export (full account data as JSON)
- Delete account (with confirmation flow)
- Opt out of usage analytics (on/off)
- View data processing information
- API key management (Founder plan — view, create, revoke)

Settings Tab: Advanced (Founder plan only)

- API key creation and management
- Webhook URL (for build-complete notifications)
- Custom domain for Live Preview (V2)
- Pipeline log access (full routing_history and failover_log)
- Pool health dashboard (which models are currently active)
- Credit usage API (programmatic access to usage data)
- Custom Decision Memory rules via JSON import
- Priority Queue position visibility

23.7 In-Chat Controls (Per Message)

Every Thinker response message has a row of action icons beneath it.

Control	Icon	Action
Copy	Copy icon	Copy full response text to clipboard
Regenerate	Refresh icon	Re-run the same request at the same Thinking Level (costs same credits)
Thinking Level up	Arrow-up icon	Re-run at one level higher (costs difference in credits)
Share	Share icon	Generate a shareable link to this response (public read-only)
Thumbs up	■	Positive feedback sent to Anthropic (no credits)
Thumbs down	■	Negative feedback — opens feedback form, sent to Anthropic
Pin	Pin icon	Pin this response to the top of the conversation
Save to project	Folder icon	Save this response as a named note in Project Memory
Version history	Clock icon	Open version history panel (project outputs only)
Rollback to this	Undo icon	Restore the project to the state at this output (project outputs only)

23.8 Right Panel — Project Details

Slides in when a project-type pipeline is active or complete.

Project Tab	Pipeline Log Tab (Founder plan)
• Project name (editable)	• Per-agent log entries (Section 9.2)
• Created date	• routing_history trace
• Last modified date	• failover_log (which models were swapped)
• Domain used	• loop_counts per agent
• Thinking Level used	• Token cost breakdown per agent
• Tech stack summary	• Failover cost vs normal cost
• Credits spent on this project	• Total session duration
• Current status (building / complete / halted)	• Model pool health at session start
• Live Preview link (if web project)	• Export log as JSON
• Download ZIP button	
• Delete project	

Version History Tab	Files Tab
• List of versions (v1, v2, v3...)	• All files uploaded in this project

• Timestamp per version	• File name, type, size, upload date
• What changed (auto-generated diff summary)	• Re-use file in new message
• Restore to this version button	• Delete file from project context
• Compare two versions side-by-side (V2)	• Files auto-deleted on session end if not saved to project
• Max 10 versions stored (older auto-deleted)	
• Rollback note: restoring creates new version	

23.9 Onboarding Flow (New User)

Step	Screen	What Happens
1	Welcome	Brief explanation of what Thinker AI does. One sentence: 'Thinker AI thinks before it builds.'
2	Choose your plan	Free Trial / Pro / Founder cards with feature comparison. Default: Free Trial.
3	First question	Thinker asks: 'What are you here to build or figure out?' — no form, just a chat prompt.
4	Thinking Level intro	After first message, a tooltip explains the four Thinking Levels. Dismissable.
5	Credit intro	After first credit is spent, a tooltip explains Think Credits. Dismissable.
6	Domain Picker intro	After second conversation, tooltip points to the Domain Picker. Dismissable.
7	Done	No more tooltips. User is in the product. Help available via '?' button in top bar.

23.10 Empty States & Zero-Data Screens

Screen	Empty State Message
New conversation	"What do you want to build, figure out, or create today?"
Conversation history (none yet)	"Your conversations will appear here."
Version history (no versions)	"No versions saved yet — versions appear here after your first build."
Decision Memory (none set)	"No saved decisions yet. Say 'always use React' and I'll remember it."
Files tab (no files)	"No files in this project yet."
Notifications (none)	"You're all caught up."
Billing history (no invoices)	"No invoices yet — your first invoice will appear after your first billing cycle."

Thinker AI V1 Architecture Spec

Addendum 4 — Sections 24, 25, 26 & 27

API & Webhook Spec | Testing & QA Strategy | Error Codes | Analytics & Business Metrics

June 2026

This addendum completes the Thinker AI V1 Architecture Specification. Section 24 defines every API endpoint and webhook the Founder plan exposes. Section 25 defines the testing strategy for every agent and pipeline component. Section 26 defines every error code and its user-facing and developer-facing message. Section 27 defines the analytics events and business metrics that determine whether V1 is succeeding.

24. API & Webhook Specification

The Thinker AI API is available to Founder plan users only. It allows programmatic access to the pipeline, credit balance, project history, and webhook delivery of pipeline events. All API calls are made server-to-server — never from a browser client.

24.1 Authentication

All API requests must include an Authorization header with a Bearer token. API keys are created in Settings → Advanced → API key management (Section 23.6). Each key has a name, creation date, last-used date, and optional expiry.

```
Authorization: Bearer tk_live_XXXXXXXXXXXXXXXXXXXX
```

```
Key format: tk_live_ prefix for production keys
```

```
tk_test_ prefix for test/sandbox keys
```

```
Rate limit: 60 requests per minute per API key (all endpoints combined)
```

```
Burst limit: 10 requests per second
```

API keys must never be embedded in client-side code. Any key detected in a public GitHub repository will be automatically revoked by the Thinker AI security monitor (V2 feature — flag for implementation).

24.2 Base URL & Versioning

```
Production: https://api.thinker.ai/v1
```

```
Sandbox: https://sandbox.api.thinker.ai/v1
```

All responses are JSON. All timestamps are ISO 8601 UTC.

API version is in the URL path – v1 is stable for the V1 release.

Breaking changes will increment the version (v2, v3).

Non-breaking additions (new fields) do not increment the version.

24.3 Endpoints

POST /pipeline/run

Run the full Thinker AI pipeline on a prompt.

Request:

```
{
  "prompt": "Build me a landing page for a SaaS product",
  "thinking_level": "high",
  "domain": "coding",
  "language": "en",
  "webhook_url": "https://yoursite.com/hooks/thinker"
}
```

Response:

```
{
  "pipeline_id": "pip_abc123",
  "status": "running",
  "thinking_level": "high",
  "estimated_credits": 15,
  "created_at": "2026-06-27T10:00:00Z"
}
```

Note: Asynchronous. Returns immediately with pipeline_id. Result delivered via webhook (Section 24.5) or poll /pipeline/{id}.

GET /pipeline/{pipeline_id}

Get the current status and result of a pipeline run.

Request:

No body – pipeline_id in URL path.

Response:

```
{
  "pipeline_id": "pip_abc123",
  "status": "complete",
  "thinking_level": "high",
  "credits_used": 16,
  "result": {
    "summary": "Landing page built with React + Tailwind",
    "output_type": "code",
    "artifacts": [
      {"name": "index.html", "type": "file", "download_url": "..."},
      {"name": "styles.css", "type": "file", "download_url": "..."}
    ],
    "strategy_brief": {...},
    "routing_history": [...]
  },
  "created_at": "2026-06-27T10:00:00Z",
  "completed_at": "2026-06-27T10:02:34Z"
}
```

Note: Status values: queued | running | awaiting_clarification | awaiting_approval | complete | failed | halted

POST /pipeline/{pipeline_id}/respond

Send a user response to a clarification question or approval prompt mid-pipeline.

Request:

```
{
  "response_type": "clarification",
  "content": "The target audience is small business owners aged 35-55"
}
```

Response:

```
{
  "pipeline_id": "pip_abc123",
  "status": "running",
  "acknowledged_at": "2026-06-27T10:01:15Z"
}
```

Note: response_type: clarification | approval | revision | rejection

DELETE /pipeline/{pipeline_id}

Cancel a running pipeline. Credits for completed stages are still charged.

Request:

No body.

Response:

```
{
  "pipeline_id": "pip_abc123",
  "status": "cancelled",
  "credits_charged": 5,
  "cancelled_at": "2026-06-27T10:01:00Z"
}
```

Note: Only pipelines with status queued or running can be cancelled.

GET /credits/balance

Get current Think Credit balance and usage for this billing period.

Request:

No body.

Response:

```
{
  "balance": 1247,
  "plan": "founder",
  "monthly_allowance": 5000,
  "used_this_month": 3753,
  "reset_date": "2026-07-01T00:00:00Z",
  "extra_credits_purchased": 500
}
```

GET /credits/history

Get a paginated list of credit transactions.

Request:

Query params: ?limit=50&cursor;=txn_xyz&action;_type=pipeline_run

Response:

```
{
```

```
"transactions": [  
  {  
    "txn_id": "txn_xyz",  
    "action": "pipeline_run",  
    "credits": -16,  
    "pipeline_id": "pip_abc123",  
    "timestamp": "2026-06-27T10:02:34Z"  
  }  
],  
"next_cursor": "txn_abc",  
"has_more": true  
}
```

Note: action_type filter: pipeline_run | file_upload | fix | rollback | purchase | monthly_grant | failover

GET /projects

List all projects for this account.

Request:

Query params: ?limit=20&cursor;=proj_xyz&status;=complete

Response:

```
{  
  "projects": [  
    {  
      "project_id": "proj_abc",  
      "name": "SaaS Landing Page",  
      "status": "complete",  
      "domain": "coding",  
      "credits_used": 66,  
      "created_at": "2026-06-27T10:00:00Z",  
      "version_count": 3  
    }  
  ],  
  "next_cursor": "proj_xyz",  
  "has_more": false  
}
```

GET /projects/{project_id}/versions

List all saved versions of a project.

Request:

No body.

Response:

```
{  
  "project_id": "proj_abc",
```

```

"versions": [
{
"version": 3,
"description": "Added mobile responsiveness",
"credits_used": 10,
"created_at": "2026-06-27T11:00:00Z",
"download_url": "...
}
]
}

```

Note: Up to 10 versions stored per project (Section 9.6).

POST /files/upload

Upload a file for use in a pipeline run.

Request:

multipart/form-data with file field. Optional: intent field (review|analyse|edit|use_as_input)

Response:

```

{
"file_id": "file_abc123",
"name": "main.py",
"type": "code",
"language": "python",
"size_bytes": 4821,
"credits_to_process": 5,
"expires_at": "2026-06-27T18:00:00Z"
}

```

Note: File is held in temporary storage for 8 hours. Pass file_id in a /pipeline/run call to use it.

24.4 Error Response Format

All API errors return a consistent JSON structure:

```

{
"error": {
"code": "INSUFFICIENT_CREDITS",
"message": "Your credit balance (3) is insufficient for this operation (requires 10).",
"details": {
"balance": 3,
"required": 10,
"top_up_url": "https://thinker.ai/credits/buy"
},
"request_id": "req_abc123",
"docs_url": "https://docs.thinker.ai/errors/INSUFFICIENT_CREDITS"
}
}

```

```
}
```

24.5 Webhooks

Thinker AI sends webhook events to a URL you specify in /pipeline/run or in Settings → Advanced → Webhook URL (used as default for all pipelines).

Event	When Fired	Key Payload Fields
pipeline.started	Pipeline begins running	pipeline_id, thinking_level, estimated_credits
pipeline.clarification_needed	Clarification Agent needs user input	pipeline_id, questions[], status: awaiting_clarification
pipeline.stage_complete	Each pipeline stage completes	pipeline_id, stage_name, stage_number, credits_used_so_far
pipeline.approval_needed	Stage 4 approval required (Section 16)	pipeline_id, preview_url, status: awaiting_approval
pipeline.complete	Pipeline finishes successfully	pipeline_id, credits_used, artifacts[], summary
pipeline.failed	Pipeline fails unrecoverably	pipeline_id, failed_at_stage, credits_charged, error_code
pipeline.halted	Loop detection or pool exhaustion happens	pipeline_id, halt_reason, credits_charged, resume_possible
credits.low_balance	Balance drops below threshold	balance, threshold, top_up_url
credits.depleted	Balance reaches 0	balance: 0, top_up_url
failover.occurred	Model failover happened (informational)	pipeline_id, agent_name, attempt_number

Webhook Security

Every webhook request includes a Thinker-Signature header. Verify this signature before processing the payload:

```
Thinker-Signature: t=1719475200,v1=hmac_sha256_hex
```

Verification (Node.js example):

```
const payload = req.rawBody;
const sig = req.headers['thinker-signature'];
const [t, v1] = sig.split(',');
const timestamp = t.split('=')[1];
const expected = crypto.createHmac('sha256', WEBHOOK_SECRET)
  .update(timestamp + '.' + payload).digest('hex');
const received = v1.split('=')[1];
if (expected !== received) throw new Error('Invalid signature');
```

Webhook secret is set in Settings → Advanced → Webhook URL. Thinker AI retries failed webhook deliveries up to 5 times with exponential backoff (1m, 5m, 30m, 2h, 8h). A webhook endpoint that fails all 5 retries is automatically disabled and the founder is notified.

25. Testing & QA Strategy

Testing a multi-agent pipeline is different from testing a conventional API. Each agent is an LLM call — deterministic testing is not possible. This section defines a layered testing strategy that covers: unit-level mocking, integration testing with real models, scenario-based QA, and launch gate testing.

25.1 Testing Layers

Layer	What Is Tested	How	When Run
Unit tests	Thinker Core routing logic, credit calculations, lookup logic	Test with mocked agent responses	On every code change
Agent mock tests	Each agent's input→output contract	Agents replaced with deterministic mocks	On every individual output. Pipeline tests
Integration tests (real models)	Full pipeline produces acceptable prompts	Use standard inputs	Once a day by staging
Failover tests	Does Model Pool Manager correctly substitute models	Force model pool failures	Weekly pipeline test
Credit accuracy tests	Are credits deducted correctly	Test every action type	Automated per action type
Load tests	Does the system hold under 50 pipeline runs	Simulate 50 pipeline runs	Before each major release
Security tests	Do file upload security rules block malicious files	Upload test files with malware pattern	On every deployment
Manual QA scenarios	End-to-end user flows that cannot be automated	Human testers follow scenario scripts	Before production release

25.2 Agent Mock Contract

Every agent must have a mock that returns a valid output schema. This is the minimum contract each agent must satisfy for CI to pass:

Agent	Mock Must Return	Invalid Output (triggers retry test)
Intent Agent	{intent_type, confidence, domain, thinking_level, next_action}	next_action is null or unknown value
Clarification Agent	{complete: bool, questions: [], requirements: {}, next_action}	complete: true with empty requirements
Strategy Agent	{goal_clarity, mvp_scope, key_risks: [], success_criteria: [], next_action}	next_action is null (Strategy cannot replan)
Planner Agent	{steps: [{id, description, output_type, dependencies}], next_action}	steps is an empty plan
Research Agent	{findings: [], sources: [], next_action}	Any malformed source URL
Builder Agent	{artifact_type, content, files: [], next_action}	Empty content string
Reviewer Agent	{passed: bool, issues: [{severity, description}], next_action}	passed: true with high-severity issues
Critic Agent	{concerns: [], severity_overall, next_action}	severity_overall not in [low, medium, high, critical]
Judge Agent	{scores: {}, approved: bool, next_action}	approved: true with any score below 40%
Consensus Agent	{votes: [{model, verdict, reason}], final_verdict, next_action}	votes count not equal to pool size
Design Agent	{images: [{url, description}], next_action}	Empty images array on success status

Agent	Mock Must Return	Invalid Output (triggers retry test)
Document Agent	{summary, key_points: [], sections: [], next_action}	key_points: [] when document has content

25.3 Manual QA Scenario Scripts

Scenario A — Simple chat

- Open app, type 'What is a REST API?'
- Assert: ■ Low Thinking selected automatically
- Assert: 1 credit deducted
- Assert: No clarification questions asked
- Assert: Direct answer returned in under 5 seconds

Scenario B — Project build (Happy path)

- Type 'Build me a todo app in React'
- Assert: ■ High Thinking selected
- Assert: Signature Question asked before pipeline runs
- Assert: Goal Discovery questions appear
- Assert: Strategy brief shown (Stage 1)
- Assert: Blueprint shown (Stage 2) — user must approve
- Assert: Build progress visible (Stage 3)
- Assert: First screen approval prompt appears
- Assert: Final ZIP downloadable (Stage 5)
- Assert: Total credits match estimate shown before run

Scenario C — Vague request triggers Goal Discovery

- Type 'I want to build something'
- Assert: Intent confidence below 70%
- Assert: Goal Discovery Mode activates
- Assert: Questions asked one at a time (not all at once)
- Assert: Pipeline does not reach Planner until discovery complete

Scenario D — Failover is invisible

- In test environment: block primary model for Builder Agent
- Submit a project build request
- Assert: Build completes successfully
- Assert: No error message shown to user
- Assert: failover_log has one entry in Pipeline State
- Assert: User sees no delay beyond 10 seconds total

Scenario E — Credit gate

- Set test account balance to 3 credits
- Submit a High Thinking request (requires 10 credits base)
- Assert: Credit confirmation modal appears
- Assert: Pipeline does not start without confirmation
- Assert: 'Insufficient credits' message shown if user confirms
- Assert: Top-up CTA visible

Scenario F — File upload (code review)

- Upload a Python file with a known bug
- Assert: File received card shown with language detected
- Assert: Review starts automatically (5 credits charged)
- Assert: Bug identified in Reviewer findings
- Assert: 'Fix this bug' option presented
- Assert: Builder fixes bug, new file downloadable

Scenario G — RTL language

- Write a message in Arabic
- Assert: Language detected as Arabic
- Assert: All agent responses in Arabic
- Assert: UI text right-aligned
- Assert: Generated code stays in English
- Assert: README in Arabic

Scenario H — Loop detection

- In test environment: force Reviewer to always return next_action: retry
- Submit any build request
- Assert: Builder retried exactly 3 times
- Assert: On 4th attempt, pipeline halts
- Assert: User notified with clear message
- Assert: Pipeline State saved with resume_from_agent set

25.4 Integration Test Prompt Set

These 10 prompts are run against the real pipeline on every staging deploy. Each must score $\geq 70\%$ on the Judge Agent's five-point rubric to pass.

#	Prompt	Expected Thinking Level	Pass Criteria
1	What is the difference between SQL and NoSQL?	Low	Accurate comparison, < 5s, 1 credit
2	Compare React vs Vue for a small team building a sales dashboard	Medium	Asks team size + app type, structured comparison

#	Prompt	Expected Thinking Level	Pass Criteria
3	Build a landing page for a fitness app	■ High	Full 5-stage pipeline, downloadable HTML output
4	Build a REST API for a blog with auth	■ High	Strategy brief mentions auth risk, Consensus fires on auth
5	I want to start a startup	■ High	Goal Discovery activates, 5-layer discovery runs
6	Validate my idea: an app that connects local farmers to restaurants (Fooder Plan)	■ High	Goal Discovery activates, multi-model assessment
7	Review this code (upload: buggy Python file)	■ Medium	Bug found, fix offered, 5 credits charged
8	Summarise this document (upload: 10-page PDF)	■ Medium	Document Agent runs, structured summary returned
9	■■■■■ ■■■■ Todo App ■■■■■■■■ ■■■■	■ High	Bengali detected, all responses in Bengali, code in English
10	Build a payment processing module	■ High	Consensus fires automatically (payment = high-risk category)

26. Error Codes & Messages

Every error in Thinker AI has a machine-readable code, a user-facing message (in the user's language, non-technical), and a developer-facing message (in English, technical). No provider or model name ever appears in any error message.

Hard rule: error messages tell the user what happened, what is saved, and what they can do next. They never expose internal implementation details.

26.1 Error Code Categories

Prefix	Category	HTTP Status Range
AUTH_	Authentication & authorisation errors	401, 403
CREDIT_	Think Credit errors	402, 429
PLAN_	Plan/feature gate errors	403
PIPELINE_	Pipeline execution errors	400, 500, 503
AGENT_	Individual agent errors	500, 503
FILE_	File upload errors	400, 413, 415
API_	API-specific errors (Founder plan)	400, 429, 500
WEBHOOK_	Webhook delivery errors	N/A (delivery-side)

26.2 Full Error Code Table

Code	HTTP	User Message	Developer Message	Action
AUTH_001	401	Please sign in to continue.	Missing or invalid Authorization header	Redirect to login
AUTH_002	401	Your session has expired. Please sign in again.	JWT token expired.	Refresh token or re-login
AUTH_003	403	You don't have permission to do that.	User does not have the required role	Show permission explanation
AUTH_004	401	This API key is invalid or has been revoked.	API key not found or revoked in database	Prompt to create new key
CREDIT_001	402	You don't have enough Think Credits to run this agent.	Credit limit reached (data not saved)	Show top-up CTA (Required).
CREDIT_002	402	Your Think Credits have run out.	Credit limit reached	Show top-up + upgrade CTA
CREDIT_003	429	You're using credits very quickly. Please wait and make limit before starting.	Request would make limit before starting	Show discussion forum
CREDIT_004	402	This project has reached its revision limit. Start a new project for this project.	Revision limit reached for this project	Show requirements for new project
PLAN_001	403	Strategy Agent is available on Pro plan.	Feature strategy_agent not available	Show upgrade CTA
PLAN_002	403	Consensus Thinking is available on Founder plan.	Feature consensus_thinking not available	Show upgrade CTA
PLAN_003	403	You've reached your monthly project limit.	Monthly project limit reached	Show upgrade CTA

Code	HTTP	User Message	Developer Message	Action
PLAN_004	403	API access is available on the Foundry plan.	API access not available on plan: {plan_name}.	Show upgrade CTA
PLAN_005	429	You already have {n} builds running on this plan. Please wait until one finishes before creating another.	Too many pipeline runs on this plan.	Show upgrade CTA or wait message
PIPELINE_001	400	I need a bit more information before I can start. Could you describe what you'd like to build?	Pipeline start rejected: describe what you'd like to build.	Redirect to build page. Clarification pending.
PIPELINE_002	500	Something went wrong during the build. Check the logs for more details.	Unhandled exception in pipeline at stage {stage_name}.	Show (stage_name) error, log error
PIPELINE_003	503	I ran into a difficulty on this step. You can either skip this step or cancel the build and try again.	All work for pipeline saved to storage. Showing build message + resume CTA	
PIPELINE_004	400	This step has been retried too many times. It may be failing because of a configuration issue.	Too many retries for stage {stage_name}. Exceeded maximum retries.	
PIPELINE_005	400	I couldn't find a pipeline with that ID.	Pipeline_id not found in database.	Return 404-style error
PIPELINE_006	400	This pipeline has already finished.	Can't cancel a completed pipeline with status {status}.	Show complete or final status
AGENT_001	503	Still working on this — taking a moment to finish up.	Modeling out after a failover. Failover initiated by {reason}.	Try to get feedback, silently! — no user action
AGENT_002	500	One of the thinking steps returned an error.	Agent expected failed during validation.	Retry up to {max_retries} attempts (PIPELINE_003)
AGENT_003	503	I wasn't able to complete this step.	All failover attempts exhausted for agent {agent_name}.	Trigger (PIPELINE_003)
FILE_001	415	That file type isn't supported. Think of it as a PDF, DOC, or XLS file.	Unsupported MIME type {mime_type}. Show supported MIME types and ZIP.	
FILE_002	413	That file is too large. The maximum file size is {max_size} MB.	File size ({size}) exceeds file type ({type}) size limit ({limit}).	compress suggestion
FILE_003	400	This file type can't be uploaded for security reasons.	Security file extension: {ext}. Matches blocked list.	no further action
FILE_004	400	This code file appears to contain sensitive information.	Sensitive data pattern detected in file {filename}. GitHub Actions can't accept code that might be sensitive.	
FILE_005	400	This code contains patterns that look like they might be harmful.	Maliciously harmful code detected in file {filename}. Show warning.	offer to analyse-only in
FILE_006	400	The ZIP file is empty or couldn't be extracted.	ZIP extraction failed: {reason}.	Ask user to re-upload
API_001	429	API rate limit reached. Please wait before making more requests.	Rate limit exceeded: {count} requests. Retries: {retries}.	Retry after header
API_002	400	The request body is missing required fields.	Missing required field: {field_name}. Return field-level validation errors	
API_003	400	The thinking_level value is not valid.	Invalid thinking_level: {value}. Must be between {min} and {max}.	Below invalid values list
API_004	403	This API key doesn't have permission for this endpoint.	API key not sufficient for endpoint {endpoint}. Return required scope	
WEBHOOK_001	N/A	N/A (delivery-side error)	Webhook delivery failed after 5 attempts.	Display endpoint URL, notify build status
WEBHOOK_002	N/A	N/A	Webhook signature verification failed.	Attempt to replay, possible replay attack or

27. Analytics & Business Metrics

This section defines what Thinker AI measures, why each metric matters, and what thresholds indicate healthy or unhealthy product behaviour. Analytics serve two purposes: product decisions (what to fix or build next) and business health (is V1 sustainable).

27.1 Analytics Principles

- Every analytics event is opt-outable by the user (Settings → Privacy → Opt out of usage analytics). Billing and security events cannot be opted out.
- No personally identifiable information (PII) is included in analytics events — user_id is a hashed pseudonym.
- No file content, prompt text, or LLM output is sent to analytics. Only structural metadata (event type, credits, duration, thinking level).
- Analytics are stored in a separate database from Pipeline State — no joins with user data are performed in the analytics layer.

27.2 Product Analytics Events

Every event below is fired from Thinker Core or the frontend. Events flow to the analytics database asynchronously — they never block the pipeline.

Event Name	Fired When	Key Properties
session_started	User opens the app	plan, language, domain_selected, is_returning_user
message_sent	User sends any message	thinking_level, has_file_attached, word_count, domain
thinking_level_selected	User manually changes Thinking Level	from_level, to_level, auto_was
pipeline_started	Full pipeline begins	thinking_level, domain, intent_type, credits_estimated
clarification_shown	Clarification questions returned to user	question_count, clarification_depth, discovery_layers_active
clarification_answered	User answers clarification	rounds_taken, goal_discovery_triggered, signature_q_answered
stage_approved	User approves Stage 1, 2, or 4	stage_number, time_to_approve_seconds
stage_rejected	User rejects or requests changes at stage	stage_number, rejection_type
pipeline_complete	Pipeline finishes successfully	thinking_level, credits_used, duration_seconds, agent_count, failure_count
pipeline_failed	Pipeline fails	failed_at_stage, error_code, credits_charged
pipeline_halted	Loop detection or pool exhaustion halt	halt_reason, agent_name, loop_count
fix_requested	User requests a fix post-delivery	fix_type (small/medium/full_rebuild), credits_used
rollback_requested	User requests version rollback	from_version, to_version, is_selective
file_uploaded	File upload completes	file_type_category, size_bytes, action_intent
image_approved	User approves a generated image	attempt_number, action_type (generate/edit/style)

Event Name	Fired When	Key Properties
image_rejected	User rejects image	attempt_number, reason_given
upgrade_prompt_shown	Upgrade CTA shown	trigger_error_code, feature_gated, current_plan
upgrade_clicked	User clicks upgrade CTA	from_plan, to_plan, trigger_feature
credits_purchased	User buys extra credits	credits_amount, price_paid, trigger (low_balance/prompted/voluntar
decision_saved	User saves a Decision Memory rule	domain, rule_type
failover_occurred	Model failover in pipeline	agent_name, primary_model_label, attempt_number, peer_audit_r
api_call_made	Founder plan API call	endpoint, method, response_time_ms, error_code (if any)

27.3 Business Health Metrics

These metrics are computed from analytics events and billing data. Target values are V1 launch targets — adjust after 30 days of real data.

Metric	Definition	V1 Target	Warning Threshold
DAU (Daily Active Users)	Unique users who send ≥ 1 message in a day	no target yet	< 10 DAU in first 2 weeks = acquisition problem
Pipeline completion rate	% of started pipelines that reach start of complete	$\geq 75\%$	$< 50\%$ = pipeline reliability problem
Clarification abandonment rate	% of users who leave during clarification	$\leq 20\%$	$> 35\%$ = too many questions or wrong questions
Free $\rightarrow$ Pro conversion rate	% of Free Trial users who upgrade to Pro in 30 days	$\geq 8\%$	$< 3\%$ = value not clear or credits too low
Pro $\rightarrow$ Founder upgrade rate	% of Pro users who upgrade to Founder within 60 days	$\geq 5\%$	$< 2\%$ = Founder plan differentiation unclear
Credit burn rate (Free Trial)	Average credits used before upgrade	30–45 or 50	< 10 = users not engaging; $= 50$ = credits too low
Monthly churn rate (Pro)	% of Pro subscribers who cancel per month	$\leq 5\%$	$> 10\%$ = product not delivering value
Average credits per pipeline (Near Thinking)	Average credits spend for a full High Thinking pipeline	40–60 credits	> 120 = cost model unsustainable
Failover rate	% of pipeline runs that trigger at least one model failover	$\leq 5\%$	$> 15\%$ = provider reliability problem
Upgrade prompt CTR	% of upgrade prompts that result in a upgrade click	$\geq 12\%$	$< 5\%$ = prompts not persuasive or mistimed
API adoption rate (Founder)	% of Founder users who make ≥ 1 API call per month	$\geq 30\%$	$< 10\%$ = API not useful or too complex
Operator LLM cost per pipeline (High Thinking API)	Average LLM cost per completed High Thinking pipeline	\$0–\$0.60	≥ 1.20 = margin unsustainable
Failover cost as % of total LLM cost	Master token spend $\div$ total token spend	$\leq 2\%$	$> 20\%$ = provider stability issue
Support ticket rate	Support tickets per 100 active users	≤ 2 per week	> 8 = UX or reliability problem

27.4 Founder Dashboard (Minimum Viable)

The founder needs a simple dashboard visible at a glance. V1 minimum: a single internal page (not user-facing) showing:

- DAU / WAU / MAU — 30-day rolling chart
- Pipeline completion rate — daily, 7-day average
- Plan breakdown — count of Free / Pro / Founder users (current)
- Credit usage today — total credits consumed across all users
- Operator LLM cost today — estimated \$ spend on model APIs
- Failover events today — count and which agents triggered
- Error rate by code — top 10 errors in last 24 hours
- New signups today / this week
- Upgrades today (Free $\rightarrow$ Pro, Pro $\rightarrow$ Founder)
- Cancellations today
- Average pipeline duration (seconds) — 7-day average

- Active pipelines right now (real-time count)

Implementation: Firestore aggregation queries or a lightweight analytics tool (e.g. Mixpanel, PostHog, or a simple Cloud Function that computes daily aggregates into a summary collection). The dashboard itself can be a simple internal React page — it does not need to be polished for V1.

27.5 V1 Success Definition

V1 is considered successful if, 60 days after launch, all of the following are true:

Criterion	Target	Why It Matters
Pipeline completion rate	≥ 70%	Core product reliability
Free → Pro conversion	≥ 8% within 30 days	Business model validation
Monthly churn (Pro)	≤ 5%	Product delivers ongoing value
Operator cost per pipeline	≤ \$0.60 (High Thinking)	Unit economics viable
Failover rate	≤ 5%	Infrastructure stable
NPS or qualitative signal	Positive user feedback visible	Product resonates

If any of these thresholds are not met at 60 days, the corresponding section of this spec should be revisited before V2 scope is planned. Metrics drive the roadmap — not assumptions.

End of Thinker AI V1 Architecture Specification — Full Edition

Sections 1–17 (Original) · Section 18 (Paid Plan) · Section 19 (Resilience) · Section 20 (Dynamic Routing) · Section 21 (Think Credits) · Section 22 (File Upload) · Section 23 (UI/UX) · Section 24 (API) · Section 25 (Testing) · Section 26 (Error Codes) · Section 27 (Analytics)